
Network Configuration with XML

NCX Overview


Andy Bierman
April 3, 2009
Netconf Central

andy@netconfcentral.com
Company Confidential

If your only tool is a hammer...



Your NM Tools




Your Tools with
NETCONF/YANG

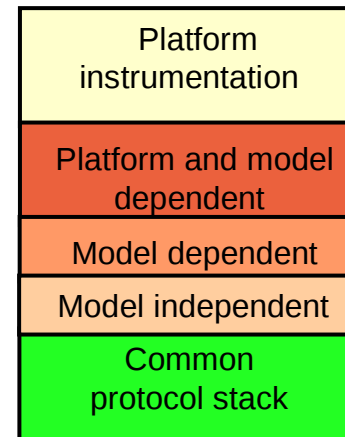
Contents

- Introduction to NCX and YANG
- NCX System Components
- NCX Implementation

NETCONF Development Costs



Without automation

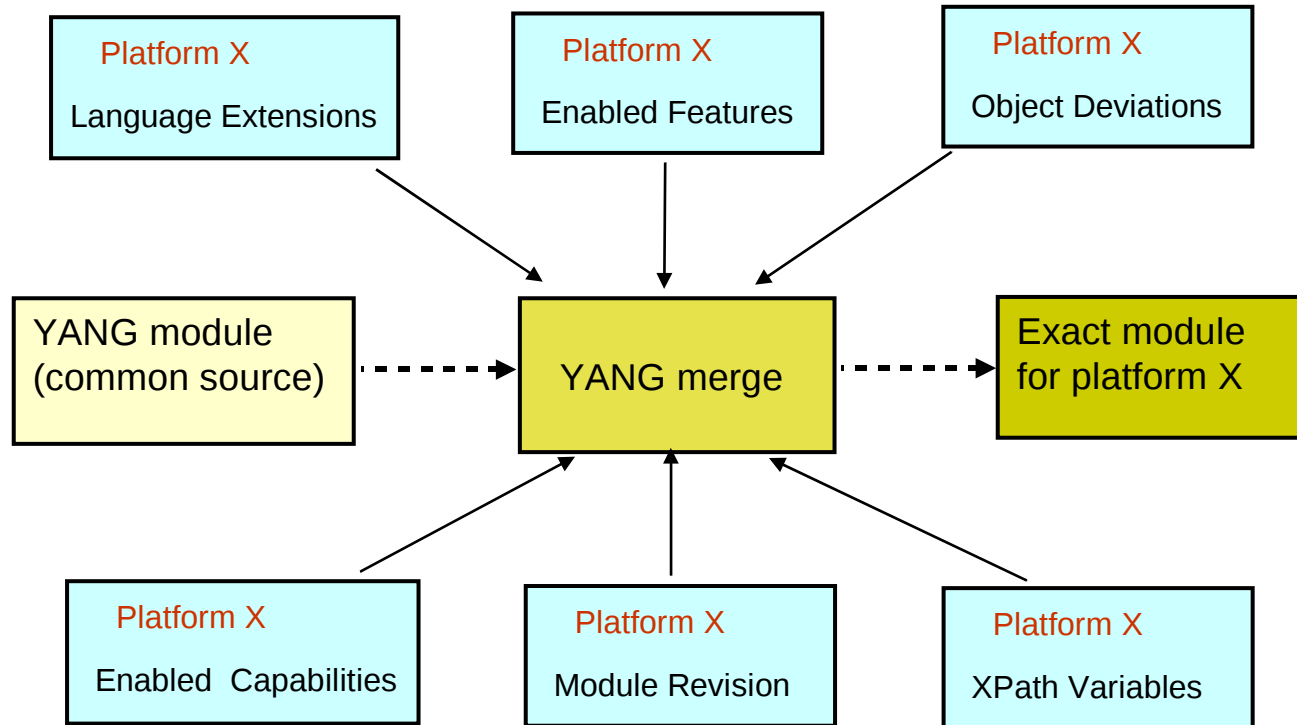


With YANG automation

What is NCX?

- Network Configuration Management Toolkit
 - » Completely Standards-based
 - NETCONF and SSH2 protocols
 - YANG data modeling language
 - XML encoding and Xpath filtering
 - ietf-netconf-state monitoring data model
 - ietf-netconf-partial-lock database locking extensions
 - » Full automation of NETCONF protocol handling
 - All session, operation, database access, error, and XML handling done in the toolkit or in the agent central stack
 - Utilization of YANG extensions, deviations, and other automation support features

YANG Cooked Modules




Platform-specific schema created automatically

What is YANG?

- Yet Another Next Generation Data Modeling Language for the NETCONF Protocol
 - » YANG will become the SMIv2 replacement for data model module definition within the Internet community
 - » Arbitrarily complex data structures are supported, making it possible for data model designers to design the MIB first, not last, saving lots of time
 - » YANG is designed from the ground up to support tool automation

YANG as Source Code



YANG-based Automation

- YANG allows clever tools to automate the development and run-time support for the entire NETCONF protocol, by using machine-readable text instead of DESCRIPTION clauses:
 - » feature/if-feature
 - » must/when statements
 - » grouping/uses/refine statements
 - » object deviations
 - » language extensions

YANG Module Partitions

- YANG features
 - » A feature defines a conceptual partition within a module
 - » Objects with the matching if-feature statement(s) are part of the logical partition
 - » Supported features are static, set at code compile-time and advertised by the agent at session start-up time
 - » A code generator can remove unneeded code for the unimplemented features on each platform

YANG Conditionals

- must and when statements
 - » A must statement defines inter-object value-space conditions, using an XPath expression
 - Each data object can have zero or more must statements
 - All the expressions must be 'true' for the object that contains the must statement(s) to be considered valid
 - » A when statement defines inter-object existence conditions, using an Xpath expression
 - Each object can have zero or one when-statement
 - The expression must be 'true' for the object that contains the when statement to be considered present in the agent
 - Objects with 'false' when expression values are automatically deleted by the agent

YANG Object Reuse

- grouping, uses, and refine statements
 - » A grouping statement defines an abstract container of reusable object definitions
 - Each grouping can contain its own types and groupings, and extend/refine other groupings via 'uses'
 - » A uses statement defines an instantiation of a particular grouping (real objects)
 - Each usage can be tailored to exact needs, using mechanisms built into YANG
 - » A refine statement 'customizes' an object used from a grouping
 - Code generators can utilize this feature as well as humans

Platform-specific Deviations

- deviation and deviate statements
 - » A deviation statement defines the alterations to one data model object for an agent implementation
 - Super-charged SNMP AGENT-CAPABILITIES
 - Built-in 'patch' operation for MIB objects
 - Contains 1 or more 'deviate' statements
 - » A deviate statement defines the patch operation
 - *not-supported* operation removes an object completely
 - *add*, *replace*, and *delete* operations manipulate the sub-clauses within an object definition
 - » Tools can automatically generate data model definitions that exactly match each agent platform

Language Extensions

- extension statement
 - » An extension statement is an annotation that defines the syntax for a tool-specific YANG language extension
 - All YANG tools must handle any extension usage gracefully, not just their own extension definitions
 - Similar to compiler directives in C code or processing instructions in XML
 - » Any complex program behavior can be automated by annotating the YANG module definitions with the right extensions
 - » Some extensions from ncx.yang:
 - root, cli, default-parm
 - secure, very-secure, xpath, hidden

Some NCX Extensions (1)

- `ncx:root;`
 - » Declares an empty container to be a database root, so it can be automatically processed and not hard-wired into RPC operations like `<get>`, `<edit-config>`, etc.
- `ncx:cli;`
 - » Declares a YANG container definition to be used as the command line and configuration file interface for any program, instead of NETCONF content
 - All documentation, compile-time, and run-time CLI code can be automatically generated with 1 extra line of YANG code
- `ncx:default-parm "foo";`
 - » Defines a default parameter name for CLI or RPC input if it is omitted by the user (used by manager applications)

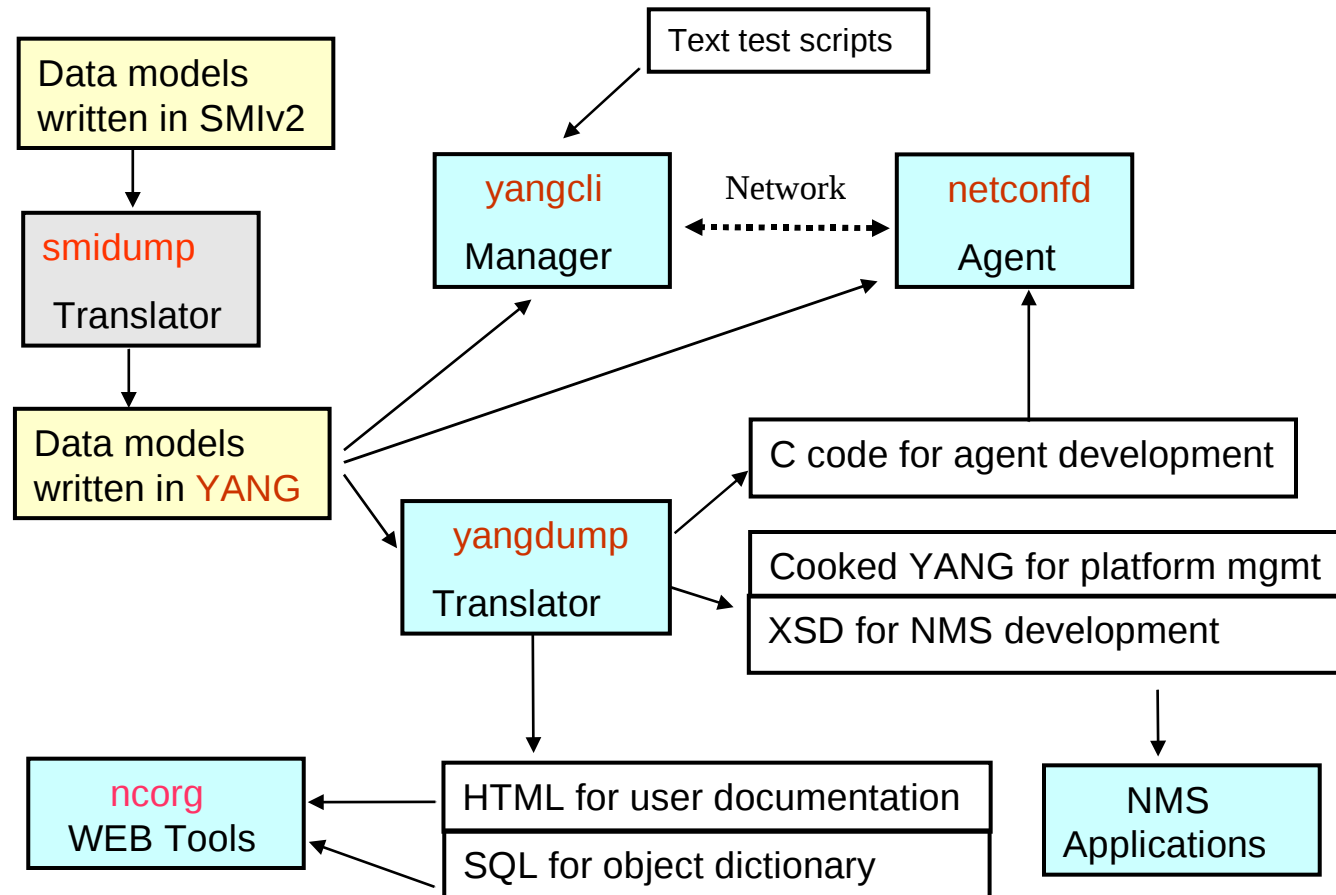
Some NCX Extensions (2)

- `ncx:secure;`
 - » Declares an object to be security sensitive
 - Only 'root' write access is allowed, even if no ACM
- `ncx:very-secure;`
 - » Declares an object to be very security sensitive
 - Only 'root' access is allowed, even if no ACM
- `ncx:xpath;`
 - » Declares an object or data type to be an Xpath expression
 - tools can automate XML namespace and validation code
- `ncx:hidden;`
 - » Removes an object from documentation and help text

Summary: Part 1

- The YANG module becomes the source code, replacing the need to manually implement almost any data model programming task
- Automation tools can be developed which remove 50% to 90% of the development costs related to network management software development
- 2nd and 3rd party application/feature development for complex networking devices is now possible

Part 2: System Components



smidump

- SMIv2 translation to YANG
 - » Part of the well-established libsmi toolset for SNMP
 - » 'smidump -f yang' converts SMIv2 to YANG
 - » A direct automated translation is needed to allow NETCONF operations (and filtering) to be applied to SMIv2-based agent instrumentation
 - » YANG could become SMIv3 in the IETF
 - Several people have already stated in IETF meetings they want to use YANG, not SMIv2 from now on
 - If all SMIv2 modules can be automatically translated with a free tool, then there is no reason to continue using SMIv2

yangdump

- Validation of YANG source files
 - » Supports all YANG constructs
 - » Full Xpath 1.0 implementation
 - » Extensive error and warning reports
 - » Full support for sub-module unification
- Translation to many output formats
 - » HTML for hyper-linked WEB docs
 - » XSD for application data model validation
 - » C code for agent instrumentation
 - » Canonical YANG for documentation and platform-specific representation of YANG files
 - » SQL for WEB application support

ncorg

- TurboGears 1 application providing the Netconf Central WEB site
 - » Uses files generated from yangdump
 - » Provides linked-docs, search, list, browse functionality for any YANG database contents
 - » Provides online validate and diff services
 - » Provides documentation and tutorial information
 - » Could be platform for additional WEB applications
 - Build a YANG module
 - User workbench front-end
 - WEB NETCONF Client application
 - Integration with bug-tracking, subversion
 - RSS feed/mailer support for module updates to database

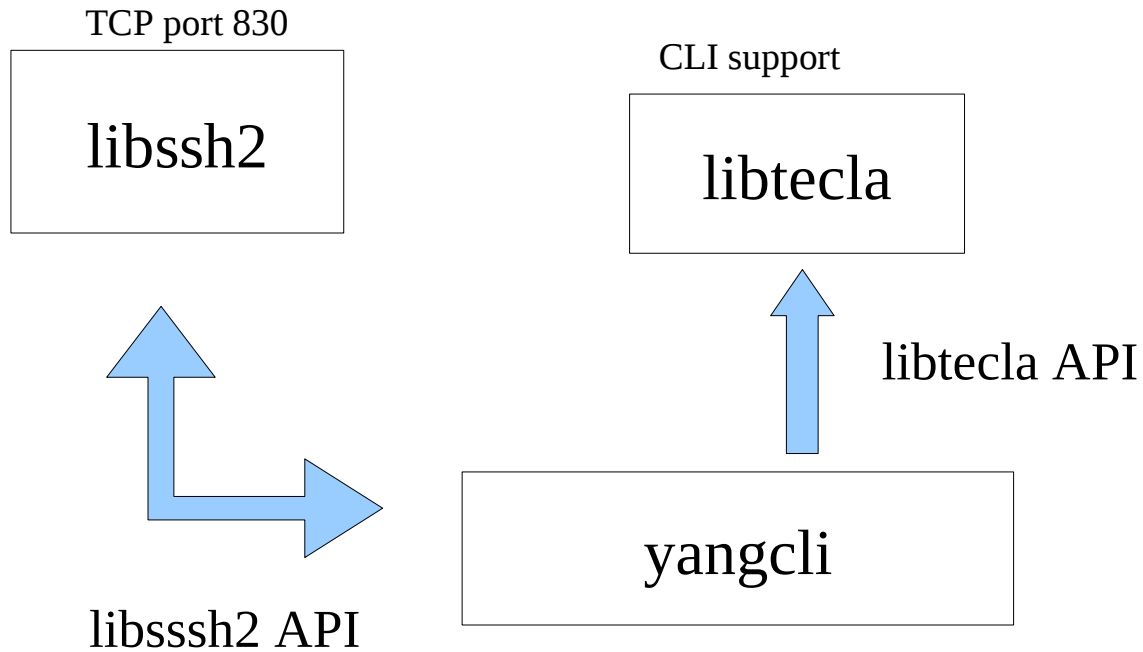
yangcli (1)

- NETCONF over SSH manager tool
 - » Full NETCONF and YANG support
 - » Completely data-driven from YANG files
 - » Context-sensitive help and commands
 - Based on the capabilities reported for each session
 - » Full Xpath 1.0 support
 - » Complex command line support via libtecla, not readline
 - » Simple script support
 - » Extensive user variable support
 - System and session parameters available to scripts
 - Output from any RPC operation can be stored in a user variable
 - `$foo = xget /interfaces/interface[name='eth0']`

yangcli (2)

- NETCONF over SSH manager tool
 - » Simple-to-use commands for common NETCONF operations:
 - 'create', 'merge', 'replace', 'delete', 'insert':
Simplified <edit-config> operations
 - 'save': agent-content specific save-changes operation
 - 'sget' and 'sget-config':
Simplified <get> and <get-config> with subtree filtering
 - 'xget' and 'xget-config':
Simplified <get> and <get-config> with XPath filtering
 - » Direct access to arbitrary YANG content
 - 'mgrload foo' is all that is needed to access all the definitions from a new module
 - Any RPC operation can be used as a CLI command


yangcli components



netconfd

- NETCONF over SSH agent
 - » Full implementation of NETCONF except:
 - :confirmed-commit is not done
 - :url is not done yet
 - » Native YANG implementation, data-driven directly from YANG source files
 - » Full implementation of all YANG-specific operations and error responses
 - » Extensive YANG-driven conf-file/CLI for customization
 - » Data models implemented:
 - ietf-netconf-state
 - ietf-partial-lock
 - ietf-with-defaults
 - nacm.yang (NETCONF Access Control Model)

netconfd components



netconfd Automation

- 100% handled in the central stack:
 - » All session, capability, XML, and YANG mechanisms
 - » All validation of every constraint available in YANG
 - » All validation of NCX extensions
 - » All aspects of the <edit-config> operation
 - » All YANG extensions, such as 'insert' and 'key'
 - » All validation and invocation of all NETCONF operations
 - » All streamed subtree and Xpath filtered output
 - » All <rpc-error> collection and reporting
 - » All <candidate/> database <commit> validation
 - » <candidate/> emulation for <running/> as target
 - » All 'false when statement' removed during commit

YANG-o-mation

- YANG is the engine that drives the bus
 - » The standard YANG source file is also the implementation source file
 - » There is no code to write
 - » The YANG file is the code
- Why automation is critical
 - » Lowers development costs by 50% to 90%
 - » Maintain consistency across platforms and releases
 - » Memory trade-off not significant because data-driven code replaces a lot of static code
 - » Domain-specific experts can focus on writing a good data model, and not deal with 'protocol weirdness'


Summary: Part 2

- The NCX tool-set includes tools to automate most aspects of NETCONF application and agent software development
- An extensible layered design allows any significant component to be replaced, for easier platform integration
- 100% standards-based solution adopted by the IETF insures multi-vendor support and the investment in data model definition modules

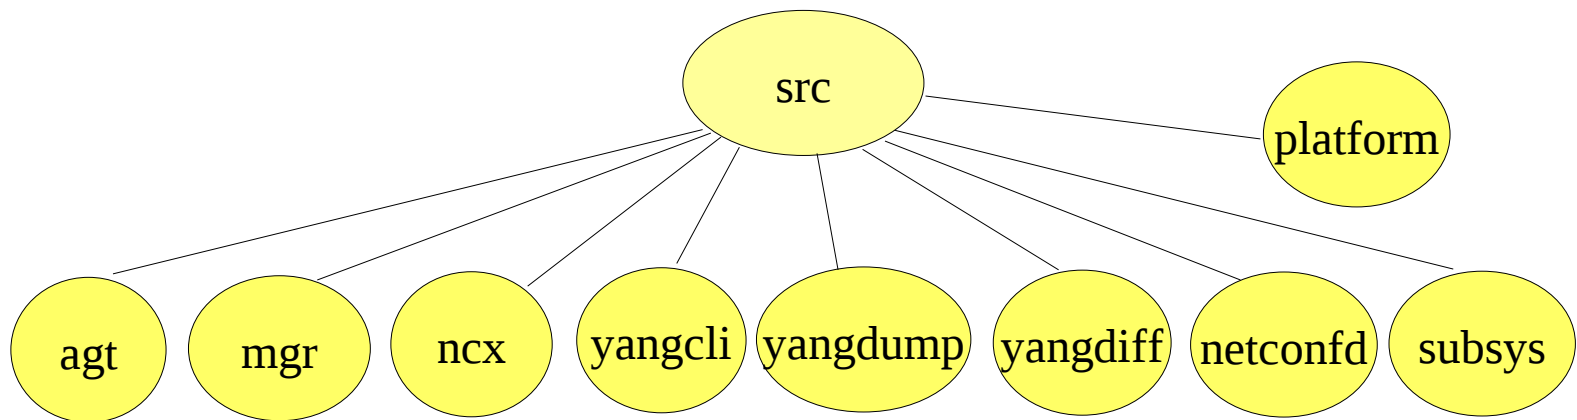
Part 3: NCX Implementation

- Approx. 180,000 lines ISO standard C code
 - » Compiled as 'std=gnu99' to support XPath floating point operations
 - » Supported platforms at this time
 - Fedora, MacOSX, Ubuntu, OpenSuse
 - Should compile on any *nix platform, but automake support is not done yet (uses proprietary procdefs.h file now)
- Approx 10,000 lines of Python and KID templates
 - » Used only in the ncorg WEB site code
- All external libraries used have acceptable licensing (e.g., BSD or MIT)
 - » glibc, libtecla, curses, libssh2, libxml2

YangTools subversion tree



C Source Tree



C Source Directories

- agt: NETCONF agent library
- mgr: NETCONF manager library
- ncx: Core common library
- yangcli: yangcli program
- yangdump: yangdump program
- yangdiff: yangdiff program
- netconfd: netconfd program
- subsys: netconf-subsystem program
- platform: platform.profile and procdefs.h

C Coding Conventions

- Super strict coding conventions
- Self-compiling H files
- No complex C coding constructs
- Braces used in every block
- Every function fully commented
- Very few `#ifdefs` in the actual code (DEBUG)

GNU Compiler Flags

- Super strict compiler checks used:

CWARN=-Wall -Wno-long-long -Wformat-y2k -Winit-self \
-Wmissing-include-dirs -Wswitch-default -Wunused-parameter \
-Wextra -Wundef -Wshadow -Wpointer-arith \
-Wwrite-strings -Wbad-function-cast -Wcast-qual -Wcast-align \
-Waggregate-return -Wstrict-prototypes -Wold-style-definition \
-Wmissing-prototypes -Wmissing-declarations \
-Wpacked -Wunreachable-code -Winvalid-pch \
-Wredundant-decls -Wnested-externs -Winline -std=gnu99 -Werror

About the Author (1)

- Computer Science degree (SFSU 1987)
- 20 years experience and almost 2 million lines of code in commercial embedded network products
 - » DAVID Systems (5 years)
Operating systems, SNMP stack and agent
 - » SynOptics/Bay Networks (3 years)
Designed and implemented RMON agent, portable agent architecture
 - » Cisco Systems (10 years)
Designed and implemented RMON blades, Entity MIB, many other SNMP agent components in IOS and CatOS
 - » Netconf Central (4 years)
Founded, designed and implemented YangTools

About the Author (2)

- 18 years leading and contributing to Network Management standards efforts in the IETF
 - » RMONMIB WG Chair and co-author (12 years)
RMON2-MIB, DSMON-MIB, APM-MIB, TPM-MIB
 - » Entity MIB WG Author (5 years)
ENTITY-MIB (1, 2, 3), ENTITY-SENSOR-MIB
 - » PSAMP (Packet Sampling) WG co-Chair (4 years)
PSAMP Architecture, Sampling, Information Model
 - » NETCONF WG Co-Chair (5 years)
NETCONF protocol and transport mappings
 - » NETMOD WG (2 years)
Major contributor to YANG documents
Author of YANG Usage Guidelines

Diagnostic Code

- SET_ERROR macro
 - » Used to generate programming exceptions for debugging
 - » All input parameters checked with SET_ERROR for all external functions (if DEBUG flag is set in make)
- Memory leak detection
 - » Internal memory trace code makes sure no memory leaks are left in the code
 - » mtrace functionality built-in (if MEMCHECK flag set in make)
- Centralized error logging facility with printf-style access functions
 - » 6 levels (error, warn, info, debug, debug2, debug3)

YANG Database

Module component templates

ncxtypes

yang_parse

obj

Object tree

Configuration database

cfg

agt_ncx_cfg_save

val

Value tree

agt_val_parse

mgr_val_parse

cli_parse

ses

Session Control Block (ses_cb_t)

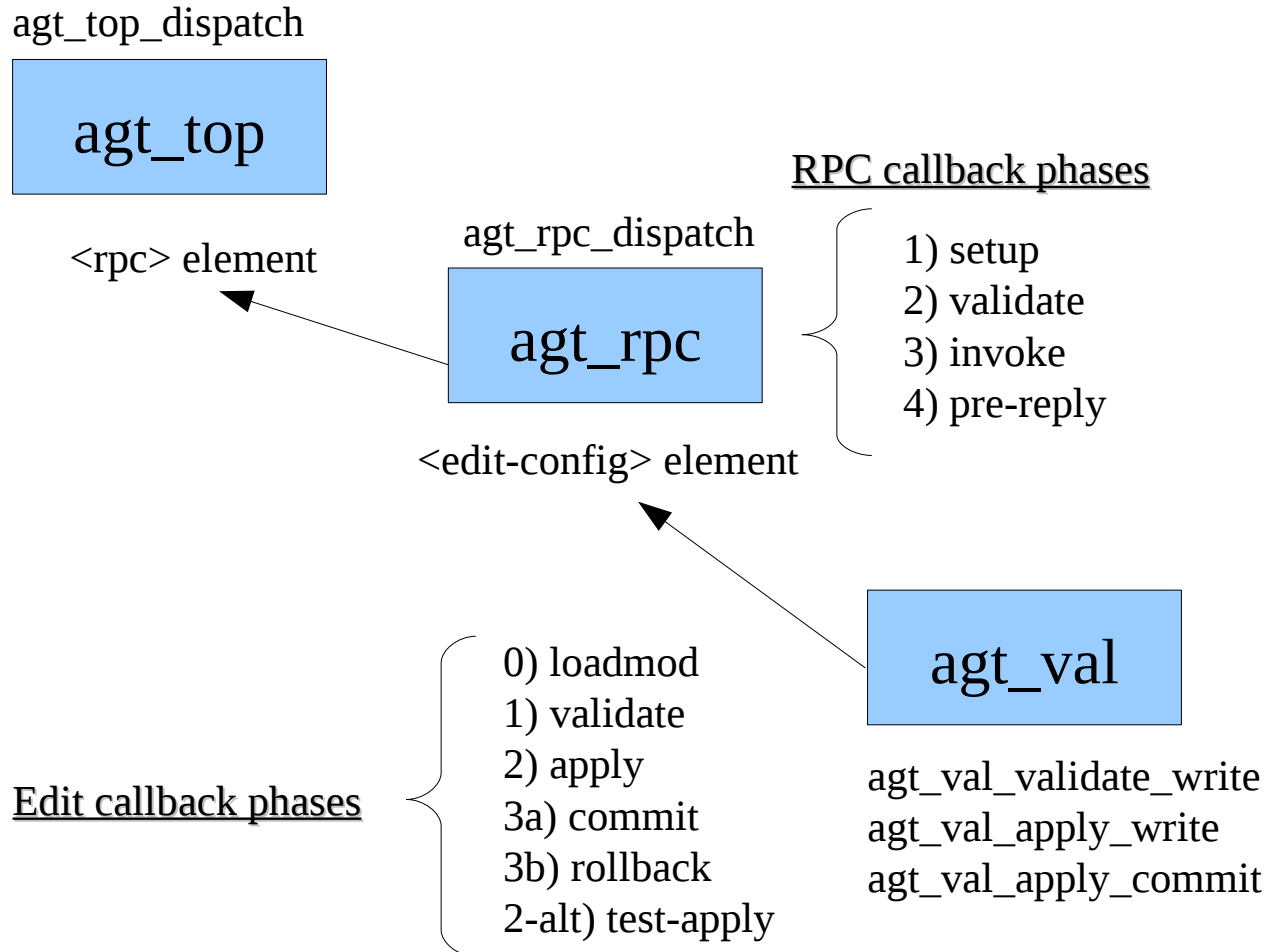
NCX Core Library

- Supports all common data structures
- YANG parsing, convert to internal object tree
- Support for all YANG datatypes and constructs
- YANG module library/revision handler
- XPath 1.0 implementation
- Hash table definition registry
- Double linked queues
- CLI, conf-file, and help text handling
- XML support and XMLNS registry
- Common error handling, logging and diagnostics

netconf/src/ncx files

b64.c	dlq.h	obj.c	ses_msg.h	xml_msg.c	yangconst.h
b64.h	ext.c	obj.h	status.c	xml_msg.h	yang_ext.c
blob.c	ext.h	obj_help.c	status.h	xmlns.c	yang_ext.h
blob.h	getcb.h	obj_help.h	tk.c	xmlns.h	yang_grp.c
bobhash.c	grp.c	op.c	tk.h	xml_util.c	yang_grp.h
bobhash.h	grp.h	op.h	top.c	xml_util.h	yang.h
cap.c	help.c	rpc.c	top.h	xml_val.c	yang_obj.c
cap.h	help.h	rpc_err.c	tstamp.c	xml_val.h	yang_obj.h
cfg.c	log.c	rpc_err.h	tstamp.h	xml_wr.c	yang_parse.c
cfg.h	log.h	rpc.h	typ.c	xml_wr.h	yang_parse.h
cli.c	Makefile	runstack.c	typ.h	xpath1.c	yang_typ.c
cli.h	ncx.c	runstack.h	val.c	xpath1.h	yang_typ.h
conf.c	ncxconst.h	send_buff.c	val.h	xpath.c	
conf.h	ncx.h	send_buff.h	val_util.c	xpath.h	
def_reg.c	ncxmod.c	ses.c	val_util.h	xpath_yang.c	
def_reg.h	ncxmod.h	ses.h	var.c	xpath_yang.h	
dlq.c	ncxtypes.h	ses_msg.c	var.h	yang.c	

Agent Callback Structure



Agent Core Library

- Provides all NETCONF agent functionality
- Key modules:
 - » agt: common init and cleanup
 - » agt_cap/agt_hello: agent capabilities handling
 - » agt_ncx: NETCONF RPC operations
 - » agt_connect/agt_ncxserver/agt_ses: session handling
 - » agt_rpc: <rpc> and <rpc-reply> handler
 - » agt_val/agt_val_parse/agt_xml: content handling
 - » agt_tree/agt_xpath: subtree and Xpath filtering
 - » agt_cb: <edit-config> callback handling
 - » agt_state: ietf-netconf-state module implementation
 - » agt_acm: nacm access control module implementation

netconf/src/agt files

agt_acm.c	agt_connect.c	agt_rpc.c	agt_state.h	agt_val.c
agt_acm.h	agt_connect.h	agt_rpcerr.c	agt_timer.c	agt_val.h
agt.c	agt.h	agt_rpcerr.h	agt_timer.h	agt_val_parse.c
agt_cap.c	agt_hello.c	agt_rpc.h	agt_top.c	agt_val_parse.h
agt_cap.h	agt_hello.h	agt_ses.c	agt_top.h	agt_xml.c
agt_cb.c	agt_ncx.c	agt_ses.h	agt_tree.c	agt_xml.h
agt_cb.h	agt_ncx.h	agt_signal.c	agt_tree.h	agt_xpath.c
agt_cli.c	agt_ncxserver.c	agt_signal.h	agt_util.c	agt_xpath.h
agt_cli.h	agt_ncxserver.h	agt_state.c	agt_util.h	Makefile

Agent Code Generation

- `yangdump --format=h --module=foo`

- » Generates constants for features and identifiers

```
#define ipfix_psamp_F_psampSampRandOutOfN 1
#define ipfix_psamp_N_exportingProcess (const xmlChar *)"exportingProcess"
#define ipfix_psamp_N_template (const xmlChar *)"template"
#define ipfix_psamp_N_observationDomainId (const xmlChar *)"observationDomainId"
```

- » Generates C equivalent typedefs for all data objects

```
/* list /ipfix/exportingProcess */
typedef struct ipfix_psamp_T_exportingProcess_ {
    dlq_hdr_t qhdr;
    xmlChar *name;
    uint32 exportingProcessId;
    dlq_hdr_t destination;
    dlq_hdr_t fileWriter;
    dlq_hdr_t options;
    dlq_hdr_t template;
    dlq_hdr_t transportSession;
} ipfix_psamp_T_exportingProcess;
```

Manager Core Library

- Provides all NETCONF manager functionality
- Key modules:
 - » mgr: common init and cleanup
 - » mgr_cap/mgr_hello: manger capabilities handling
 - » mgr_io/mgr_ses: session handling
 - » mgr_top: top level element registry
 - » mgr_rpc: <rpc> and <rpc-reply> handler
 - » mgr_val_parse/mgr_xml: content handling

netconf/src/mgr files

Makefile	mgr.h	mgr_io.h	mgr_ses.h	mgr_top.h	mgr_xml.h
mgr.c	mgr_hello.c	mgr_rpc.c	mgr_signal.c	mgr_val_parse.c	
mgr_cap.c	mgr_hello.h	mgr_rpc.h	mgr_signal.h	mgr_val_parse.h	
mgr_cap.h	mgr_io.c	mgr_ses.c	mgr_top.c	mgr_xml.c	

Summary: Part 3

- The NCX NETCONF/YANG implementation is complete and ready to be adapted to many platforms
- Security, protocol conformance, and automation are built-in features from the ground up
- Many content-based add-on features are possible, which would continue to increase toolkit value over time